

Elastic Load Balancing (ELB)

Content Prepared By: Chandra Lingam, Cotton Cola Designs LLC

For Distribution With AWS Certification Course Only

Copyright © 2017 Cotton Cola Designs LLC. All Rights Reserved.

All other registered trademarks and/or copyright material are of their respective owners



Elastic Load Balancing

- Automatically distribute traffic across EC2 instances
- Health Check – route traffic to healthy instances
- Achieve fault tolerance protect against:
 - EC2 instance failure
 - Availability Zone failure
- Fully managed – automatically scales load balancing capacity needed to route application traffic
- Single point-of-contact for your clients

Elastic Load Balancing

- Add or Remove EC2 instances without impacting clients
- Protection against Denial-Of-Service Attacks – absorb certain type of DOS attacks
- Reduce attack surface – for internet facing services, keep your EC2 instances in private subnet
- Use as an Internal Load Balancer or Internet Facing Load Balancer
- SSL Offload

Elastic Load Balancer - Concepts

[Internet Facing Load Balancer](#)

[Internal Load Balancer](#)

Load Balancer Types

- Classic Load Balancer
 - Simple load balancing across EC2 instances
 - Supports HTTP, HTTPS, TCP, SSL (Secure TCP)
 - OSI Layer 4 (Transport), 7 (Application) Load Balancer
- Application Load Balancer
 - Path based routing
 - Route traffic to multiple services
 - Route traffic to different ports on the same EC2 instance
 - Ideal for microservices and container based architectures
 - Supports HTTP, HTTPS, HTTP/2, WebSocket
 - OSI Layer 7 (Application) Load Balancer

WebSocket

- [WebSocket](#) provides Full Duplex communication over a single TCP connection
- Facilitates near real-time data transfer by sending data back and forth keeping the connection open
- Ideal for frequent interaction between Web browser and Web Server

HTTP/2

- [HTTP/2](#) (HTTP version 2) is an improvement over HTTP protocol
- More efficient use of network resources
- Reduced perception of latency – multiple concurrent exchanges on same connection
- Server to Client Push notification support – Web server can send newer information to client without requiring the client to request for data [[‘Note that you can't use the server-push feature of HTTP/2’ in ELB](#)]

Cross Zone Load Balancing

- For fault tolerance, EC2 instances should be distributed across two or more availability zones
- Cross Zone Load Balancing controls how traffic is distributed across Availability Zones and Instances in each Availability Zone

Disabled - Cross Zone Load Balancing

- Distribute traffic evenly across Availability Zones
 - Happens when Cross Zone load balancing is “Disabled”
 - Two availability zones ‘A’ and ‘B’ would each receive 50% of the traffic irrespective of number of EC2 instances in each Availability Zone
 - May cause higher loading if one Availability Zone has fewer EC2 instances
 - Default mode in classic load balancer

Enabled - Cross Zone Load Balancing

- Distribute traffic evenly across EC2 instances in all Availability zones
 - Happens when Cross Zone load balancing is “Enabled”
 - Availability Zone ‘A’ has 3 instances and Availability Zone ‘B’ has 2 instances. Each instance would receive 20% of the traffic (1/5th)
 - Default mode in application load balancer

Internet Facing Load Balancer

- One public subnet in every enabled Availability Zone
- Load Balancer nodes are deployed in the public subnet with public IP addresses
- Public DNS host name resolves to public IP Addresses
- Ensure subnet has at least 8 IP addresses left for load balancer use
- Your EC2 instances can be deployed in private subnet
- Load Balancer talks to your EC2 instances using Private IP Addresses

Internal Load Balancer

- One private or public subnet in every enabled Availability Zone
- Load Balancer nodes are deployed in the specified subnet with private IP addresses
- Public DNS host name resolved to private IP Addresses
- Ensure subnet has at least 8 IP addresses left for load balancer use
- Your EC2 instances can be deployed in private subnet
- Load Balancer talks to your EC2 instances using Private IP Addresses

Classic Load Balancer

- Create Classic Load Balancer
- Specify Listener Configuration
 - Frontend - Protocol and Port where load balancer listens for connections from client
 - Backend - Protocol and Port where EC2 instances listen. Load balancer forwards the requests to the Backend
- Specify Subnets - One Subnet per Availability Zone for Load Balancer nodes
- Assign Security Groups for Load Balancer
- Configure Health Checks
- Register EC2 Instances

Classic Load Balancer Demo

Internet Facing Load Balancer

Internal Load Balancer

Classic Load Balancer Demo

- Internet Facing Load Balancer
- VPC with 2 Availability Zones
- Each AZ contains one public and one private subnet
- Load Balancer nodes in public subnet
- EC2 instances in private subnet
- Bastion Host in public subnet (maintenance access to webserver instances)
- Health Check and simulate errors

Classic Load Balancer

Listener Configurations

Monitoring

[CloudWatch Metrics](#)

[Access Logs](#)

[CloudTrail Logs](#)

Application Load Balancer

- [Application Load Balancer](#) – Single Point of contact for clients
- Listener Configuration
 - Frontend - Protocol and Port where load balancer listens for connections from client
 - Rules – Routing rules based on condition, priority and target group. When condition is met, request is forwarded to the specified Target Group
 - Backend / Target Group – Each Target Group routes request to registered EC2 instances at specified protocol and port
 - EC2 Instance – Can be registered many target groups

Application Load Balancer

- Health Checks – Configured for every Target Group
 - Health checks are performed on every registered instance in the target group
- Specify Subnets - One Subnet per Availability Zone for Load Balancer nodes
- Assign Security Groups for Load Balancer

Application Load Balancer Benefits

- Path based routing – route based on URL
 - Microservices
 - Containerized application
- Route to multiple services on a Single EC2 instance by registering with different target groups and ports
- CloudWatch metrics at Target Group level
- Improved load balancer performance

Elastic Load Balancer Types

[Table: Features of Classic and Application Load Balancer](#)

Routing Algorithm

Classic Load Balancer

- HTTP/HTTPS Listener - Least Outstanding Requests routing
- TCP Listener - Round Robin

Application Load Balancer

- Round Robin within Target Group

Billing

Classic Load Balancer Billing

- Hourly Charge for load balancer (\$0.025 per hour)
- Per GB data processed charge (\$0.008 per GB)

Application Load Balancer Billing

- Hourly Charge for load balancer (\$0.0225 per hour)
- Hourly Load Balancer Compute Unit (LCU) (\$0.008 per LCU hour)